

ROSE: A Reward-Oriented Data Selection Framework for LLM Task-Specific Instruction Tuning

Yang Wu^{♣*†} Huayi Zhang^{♡*} Yizheng Jiao[♡] Lin Ma[♡]
Xiaozhong Liu[♣] Jinhong Yu[♣] Dongyu Zhang[♡] Dezhi Yu[♡] Wei Xu[♡]
♣Worcester Polytechnic Institute ♡ByteDance Inc.
{ywul19, xliu14, jyu7}@wpi.edu
{huayi.zhang, yizhengjiao, lin.mal}@bytedance.com
{dongyu.zhang, dezhi.yu, wei.xu}@bytedance.com

Abstract

Instruction tuning has underscored the significant potential of large language models (LLMs) in producing more human controllable and effective outputs in various domains. In this work, we focus on the data selection problem for task-specific instruction tuning of LLMs. Prevailing methods primarily rely on the crafted similarity metrics to select training data that aligns with the test data distribution. The goal is to minimize instruction tuning loss on the test data, ultimately improving performance on the target task. However, it has been widely observed that instruction tuning loss (i.e., cross-entropy loss for next token prediction) in LLMs often fails to exhibit a monotonic relationship with actual task performance. This misalignment undermines the effectiveness of current data selection methods for task-specific instruction tuning. To address this issue, we introduce ROSE, a novel **R**eward-Oriented **i**n**S**truction data **s**election method which leverages pairwise preference loss as a reward signal to optimize data selection for task-specific instruction tuning. Specifically, ROSE adapts an influence formulation to approximate the influence of training data points relative to a few-shot preference validation set to select the most task-related training data points. Experimental results show that by selecting just 5% of the training data using ROSE, our approach can achieve competitive results compared to fine-tuning with the full training dataset, and it surpasses other state-of-the-art data selection methods for task-specific instruction tuning. Our qualitative analysis further confirms the robust generalizability of our method across multiple benchmark datasets and diverse model architectures.¹




Figure 1: Illustration of our Reward-Oriented Instruction Data Selection (ROSE) approach compared with the traditional method. Unlike the traditional focus on minimizing validation loss, ROSE maximizes task-specific reward to more precisely align with real-world task performance.

1 Introduction

While large language models (LLMs) are widely recognized for their strong generalization capabilities, many fields require enhanced domain-specific performance, e.g., health monitoring (Kim et al., 2024b), legal question answering (Wu et al., 2024), and mathematics tutoring (Li et al., 2023b). Instruction tuning has emerged as a popular method for adapting foundation models to specialized tasks, which typically involves curating a high-quality training dataset. Although recent advancements in open-source datasets and synthetic data generation have facilitated the generation of large training datasets, it is widely acknowledged that the quality of training data is more crucial than its quantity in instruction tuning (Chen and Mueller, 2024; Xia et al., 2024). Consequently, practitioners must carefully select high-quality data to enhance the model’s capabilities for specific tasks. This challenge is further compounded by the complexity of domain-specific requirements and the black-box properties of LLMs, rendering it nearly infeasible for humans to manually select the most suitable training set. Therefore, developing more effective data selection methods is becoming increasingly crucial for reducing training costs and efficiently optimizing instruction tuning for specific tasks.

Despite various methods proposed for instruction tuning data selection (Du et al., 2023; Mekala et al., 2024; Agarwal et al., 2024), identifying high-quality

*Equal contribution. Initial work was conducted while Y.W. was an intern at ByteDance.

†Corresponding author.

¹<https://github.com/YANGWU001/ROSE.git>

instruction tuning data for target tasks remains a significant challenge. Existing methods typically design specific similarity metrics to select a set of candidates samples whose distribution aligns with that of the target task data. For instance, DSIR (Xie et al., 2023) uses n-gram feature similarity to enhance the selection process of relevant data samples. Moreover, RDS (Zhang et al., 2018) and LESS (Xia et al., 2024) calculate the similarity on the model embedding and gradient space to capture task-specific semantics and model characteristics, respectively. Despite their successes, these similarity based methods are fundamentally limited.

Our analysis reveals that these strategies hinge on Empirical Risk Minimization (ERM), selecting training data that mirrors the target task data distribution by minimizing training loss, particularly next-token prediction loss. However, it is widely acknowledged that next-token prediction loss often fails to accurately reflect a model’s real-world performance (e.g., alignment degree with human preference, reasoning ability on complex math problem) on target task (Zhou et al., 2024; Tay et al., 2021). This significant gap between the implicit theoretical foundation of similarity-based methods and practical task performance limits the effectiveness of these methods in instruction tuning for task-specific fine-tuning. In response to these insights, we introduce a novel reward-oriented instruction data selection (ROSE) method in Figure 1, which shifts the intrinsic selection objective from minimizing validation cross entropy loss to maximize the reward for the target task. Inspired by Direct Preference Optimization (DPO) (Rafailov et al., 2024), our approach utilizes a few-shot set of pairwise samples as the task-specific preference validation set, which we assume reflects the desired LLM’s performance on the target task, and we use the DPO loss function to approximate the expected reward value of the trained LLM on the preference validation data. By leveraging the gradient-based influence estimation techniques, ROSE is able to select the instruction tuning samples that leads a downstream LLM with optimized performance on the target task.

We conduct comprehensive experiments across various datasets and model architectures. Experimental results show that our method outperforms existing similarity-based techniques, including token-wise, embedding-based, and gradient-based methods. These results suggest that focusing on reward maximization, rather than loss minimization, is a promising new direction for improving task-specific fine-tuning outcomes.

The main contributions are as follows:

- **Identifying Limitations of Similarity-Based Approaches:** Our analysis reveals the limitations of the implicit theoretical foundations of the common similarity-based data selection methods. These methods focus on minimizing training loss (e.g., next-token prediction loss), which often fails to capture real-world task performance.
- **Introducing Reward-Oriented Data Selection:** We

shift the data selection objective from loss minimization to reward maximization on the target task dataset, which contains a small set of user-provided positive and negative samples that reflect the task-specific performance.

- **Propose to Approximate Reward with DPO Loss:** Our method incorporates gradient-based influence estimation techniques to select high-quality training data that optimizes task reward, improving the fine-tuning process.
- **Comprehensive Experimental Validation:** Through experiments on various datasets and models, we demonstrate that ROSE consistently outperforms existing similarity-based methods, such as token-wise, embedding-based, and gradient-based approaches, in task-specific fine-tuning.

2 Related Work

Data Selection for Instruction Tuning. Instruction tuning is crucial for aligning large language models (LLMs) with human needs (Wang et al., 2024), providing a controlled and safe method to enhance LLMs’ responsiveness and accuracy in specific domains. Wei et al. (2023) introduced InstructionGPT-4 for multi-modal large language fine-tuning, which involves encoding visual and textual data into vectors to train a trainable data selector. Furthermore, RDS (Zhang et al., 2018; Hu et al., 2023) employs the model’s last hidden layer to assess the similarity between training and validation data points, while DSIR Xie et al. (2023) uses n-gram features to assign importance weights to training samples for data selection in instruction fine-tuning. Another notable study, LESS Xia et al. (2024), follows a similar approach by selecting the most influential data from the training corpus based on the gradient similarity score of training data points with validation data points. Furthermore, NUGGETS (Li et al., 2023c) introduces a one-shot learning framework that selects high-quality instruction data by evaluating their impact on an anchor set, enabling efficient filtering of informative samples. Cherry (Li et al., 2023a) adopts a self-guided selection strategy based on Instruction-Following Difficulty (IFD), allowing models to prioritize training samples that better align with human instructions. DEITA (Liu et al., 2023) further enhances data selection by integrating complexity, quality, and diversity metrics, optimizing instruction tuning with significantly fewer training samples. While these methods improve data selection beyond naive heuristics, they rely on indirect proxies such as token perplexity or instruction difficulty, which do not establish a direct monotonic relationship with final task performance. In contrast, ROSE optimizes training data selection through pairwise preference modeling, ensuring a stronger alignment between training objectives and real-world win rate improvements.

Data Attribution and Influence Functions. The influence calculation of training data points is a piv-

total technique for detecting mislabeled samples (Deng et al., 2024; Zhang et al., 2021b, 2024; Hofmann et al., 2022), facilitating model interpretation (Madisen et al., 2022; Wu et al., 2023; VanNostrand et al., 2023), and analyzing memorization effects (Feldman and Zhang, 2020). Specifically, influence functions (Koh and Liang, 2017; Yao et al., 2025a) offer a counterfactual method to assess both model behaviors and the contributions of training data. Despite their potential, the robustness and effectiveness of these functions remain limited, particularly in the context of large language models (LLMs), where their computational demands are significant. While recent studies, such as those by Park et al. (2023), propose relatively efficient estimations of influence functions for selecting pretraining data, these methods still require complex comparisons of model training with and without the inclusion of specific data points. In line with the approach by Xia et al. (2024), we advocate that first-order influence approximations are effective for data selection during instruction tuning in LLM environments.

Large Language Model Alignment. LLM alignment aims to train large language models (LLMs) to behave in ways that align with human expectations. A primary approach for this is Reinforcement Learning from Human Feedback (RLHF), which tunes LLMs to reflect human preferences and values (Ziegler et al., 2019; Bai et al., 2022). It has been effectively applied in various domains, including enhancing model helpfulness (Tian et al., 2023), improving reasoning capabilities (Wu et al., 2025), and mitigating toxicity (Korbak et al., 2023). Despite its effectiveness, RLHF as an online preference optimization algorithm, poses significant challenges and complexities, since it is typically trained using Proximal Policy Optimization (PPO) (Schulman et al., 2017). And the high computational cost and sample inefficiency of PPO make RLHF difficult to scale effectively. In contrast, Direct Preference Optimization (DPO) (Rafailov et al., 2024) offers a simpler and more efficient offline alternative. Recent research has extended DPO beyond traditional pairwise comparisons to include evaluations across multiple instances (Liu et al., 2024; Yuan et al., 2024; Yao et al., 2025b). Further advancements have broadened preference optimization objectives, including those independent of reference models (Xu et al., 2023; Meng et al., 2024). These reward-oriented methods outperform models trained with next-token prediction loss in satisfying user preferences, as they directly optimize preference signals, whereas next-token prediction loss focuses on minimizing the difference between generated outputs and predefined labels, which may not reflect user-specific preferences.

3 Preliminaries and Background

Problem Definition. We tackle the challenge of data selection for task-specific instruction tuning. Our goal is to curate a subset $\mathcal{D}_{\text{train}}$ from a broad and compre-

hensive instruction tuning corpus \mathcal{D} , such that training a model on $\mathcal{D}_{\text{train}}$ to maximize a reward r that reflects true performance on a task-specific validation set \mathcal{D}_{val} , therefore performs well on test dataset $\mathcal{D}_{\text{test}}$. \mathcal{D}_{val} can be a few-shot dataset involving multiple target tasks, and the $\mathcal{D}_{\text{test}}$ is a fixed sample set with the same tasks in \mathcal{D}_{val} . We use Ω parametrized by θ to denote the model used for data selection and Γ parametrized by θ' to represent the final trained model.

Analysis of Similarity Based Methods. Let $\mathbb{Z} = \mathbb{R}^d$ be the d -dimensional intrinsic representation space, where the similarity-based methods calculate the similarity between the training and validation samples. Let $p_t(z)$ and $p_v(z)$ be the probability density value of the selected training set $\mathcal{D}_{\text{train}}$ and validation set \mathcal{D}_{val} , respectively, where $z \in \mathbb{Z}$. When a downstream LLM Γ is well trained on the selected training set $\mathcal{D}_{\text{train}}$, its training loss on $\mathcal{D}_{\text{train}}$ is expected to be close to 0, i.e., $\mathbb{E}_{\text{train}}[l(z, \theta')] = \int l(z, \theta') p_t(z) dz \approx 0$, where $l(z; \theta')$ represents the average of token-wise cross entropy loss in the response sequence of z .

The objective of similarity-based methods is to select a training set $\mathcal{D}_{\text{train}}$ which is independent and identically distributed (IID) with respect to the validation set (Xie et al., 2023; Zhang et al., 2021a). Under this condition, for any $z \in \mathbb{Z}$, $p_t(z) = p_v(z)$. Let $\mathbb{E}_{\text{val}}[l(z, \theta')] = \int l(z, \theta') p_v(z) dz$ be the expected loss value of Γ on the validation set. By simply substituting $p_v(z)$ with $p_t(z)$, we can infer that the downstream LLM is supposed to have small loss value of $l(z; \theta')$ on the validation set, i.e., $\mathbb{E}_{\text{val}}[l(z, \theta')] = \int l(z, \theta') p_t(z) dz \approx 0$.

Unfortunately, existing studies (Zhou et al., 2024; Xia et al., 2024; Tay et al., 2021) commonly find that the decrease in validation loss does not always lead to improved test performance in task-specific instruction tuning. Furthermore, achieving such an IID condition is often unrealistic due to the complexity of the representation space \mathbb{Z} . Therefore, it is reasonable to conclude that the intrinsic limitation of large models, namely the gap between next-token prediction loss and actual performance on downstream tasks, compromises the effectiveness of existing data selection methods for task-specific instruction tuning.

Influence Estimation Scheme. Assume the selection model LLM, denoted as Ω , assesses the influence of training data points with respect to a set of validation samples \mathcal{D}_{val} , which represents the model’s capability on specific tasks. We denote the average loss value of Ω on the validation set as $L(\mathcal{D}_{\text{val}}; \theta)$.

For simplicity, we assume the LLM is trained with a batch size of 1 using the SGD optimizer. At training step t , the contribution of a training sample z corresponds to the difference between the validation losses $L(\mathcal{D}_{\text{val}}; \theta)$ and $L(\mathcal{D}_{\text{val}}; \theta_{t-1})$, i.e., $L(\mathcal{D}_{\text{val}}; \theta) - L(\mathcal{D}_{\text{val}}; \theta_{t-1})$. Using a first-order Taylor expansion, this becomes:

$$L(\mathcal{D}_{\text{val}}; \theta_t) - L(\mathcal{D}_{\text{val}}; \theta_{t-1}) = \langle \nabla_{\theta} L(\mathcal{D}_{\text{val}}; \theta_{t-1}), \delta\theta \rangle \quad (1)$$

where $\langle \cdot, \cdot \rangle$ denotes the inner product, and $\delta\theta = \theta_t - \theta_{t-1}$ represents the change in θ at step t . With the SGD optimizer, $\delta\theta = -\alpha \cdot \nabla_{\theta} L(z; \theta_{t-1})$, where α is the learning rate and $\nabla_{\theta} L(z; \theta_{t-1})$ denotes the gradients of the loss with respect to the training sample z . Substituting this into the equation, we get:

$$L(\mathcal{D}_{\text{val}}; \theta_t) - L(\mathcal{D}_{\text{val}}; \theta_{t-1}) \propto \langle \nabla_{\theta} L(\mathcal{D}_{\text{val}}; \theta_{t-1}), \nabla_{\theta} L(z; \theta_{t-1}) \rangle \quad (2)$$

Eq. 2 shows that the inner product between the gradient of the loss on the training sample z and the gradient of the average loss on the validation set \mathcal{D}_{val} effectively estimates the degree to which a training sample contributes to the model’s performance. A positive gradient inner product indicates that the training sample z positively impacts the model’s performance.

Note that in the typical LLM training settings, the optimizer is usually a variant of Adam, and the training batch size is often larger than 1. This creates interactions between training samples and across batches. To address this, we adopt the heuristic-based methods used in LESS, performing a warm-up training on the LLM Ω and applying a variant of the gradient $\nabla_{\theta} L(z; \theta_{t-1})$ to reduce the discrepancy.

4 ROSE: Reward-Oriented inStruction data sElection

In this section, we introduce ROSE, a reward orientated data selection method for task-specific instruction tuning. We start from formulating the optimization objectives of ROSE, followed by a detailed explanation of implementation strategies.

4.1 Optimization Framework

Motivated by the analysis in Section 3, the objective of ROSE is to select a subset $\mathcal{D}_{\text{train}}$ that leads to a downstream LLM Γ with maximized reward value on validation set \mathcal{D}_{val} . Formally, we define $\mathcal{D}_{\text{val}} = \{(x^i, y^i)\}_{i=1}^{|\mathcal{D}_{\text{val}}|}$, where x, y denote the prompt and response of samples from \mathcal{D}_{val} , respectively. Inspired by Reinforcement Learning from Human Feedback (RLHF) (Bai et al., 2022) and further study Direct Preference Optimization (DPO) (Rafailov et al., 2024), we define the reward function r for ROSE through a closed-form expression as:

$$r(x, y) = \beta \log \frac{\Omega_{\theta}(y | x)}{\Omega_{\text{ref}}(y | x)} + \beta \log Z(x) \quad (3)$$

where Ω_{θ} and Ω_{ref} represent the policy and reference models, respectively. The term $\log Z(x)$ represents the partition function, and β signifies the parameter that controls the deviation from the baseline reference model². For optimization efficiency, we advance our methodology by first transforming the few-shot validation dataset \mathcal{D}_{val} into a few-shot preference validation set $\mathcal{D}'_{\text{val}} = \{(x^i, y_w^i, y_l^i)\}_{i=1}^{|\mathcal{D}'_{\text{val}}|}$. where (x, y_w, y_l) refers preference pairs from the dataset $\mathcal{D}'_{\text{val}}$, comprising the prompt, a winning response, and a losing response. By integrating this reward structure with the Bradley-Terry (BT) ranking model (Bradley and Terry, 1952), we leverage the probability of preference data directly through the policy model, formulating the following optimization objective:

$$\mathcal{L}_{\text{ROSE}}(\Omega_{\theta}; \Omega_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}'_{\text{val}}} \left[\log \sigma \left(\beta \log \frac{\Omega_{\theta}(y_w | x)}{\Omega_{\text{ref}}(y_w | x)} - \beta \log \frac{\Omega_{\theta}(y_l | x)}{\Omega_{\text{ref}}(y_l | x)} \right) \right] \quad (4)$$

where σ is the logistic function. We can implicitly define $\hat{r}_{\theta}(x, y) = \beta \log \frac{\Omega_{\theta}(y | x)}{\Omega_{\text{ref}}(y | x)}$, thus the gradient with respect to the parameters θ can be written as:

$$\nabla_{\theta} \mathcal{L}_{\text{ROSE}}(\Omega_{\theta}; \Omega_{\text{ref}}) = -\beta \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}'_{\text{val}}} \left[\sigma(\hat{r}_{\theta}(x, y_l) - \hat{r}_{\theta}(x, y_w)) \cdot \left(\nabla_{\theta} \log \Omega(y_w | x) - \nabla_{\theta} \log \Omega(y_l | x) \right) \right] \quad (5)$$

4.2 Implementations

Here, we describe how ROSE adapts the Equation 5 to select instruction data that can effectively improve model capability in target tasks, and illustrate our method in Figure 2.

4.2.1 Build Few-Shot Preference Validation Data

The current validation set used in task-specific instruction tuning typically adheres to a Supervised Fine-Tuning (SFT) format, featuring a prompt and its corresponding response. However, the widely used evaluation metric for test data in instruction tuning is the win rate, which assesses the frequency at which a target model’s response is deemed superior compared to the original test dataset response, as determined by an LLM evaluator. To align with the win rate employed in downstream tasks, we transform the few-shot SFT validation set $\mathcal{D}_{\text{val}} = \{(x^i, y^i)\}_{i=1}^{|\mathcal{D}_{\text{val}}|}$ into a preference format, denoted as $\mathcal{D}'_{\text{val}} = \{(x^i, y_w^i, y_l^i)\}_{i=1}^{|\mathcal{D}'_{\text{val}}|}$. This transformation involves generating additional responses to

²In practice, we use the pretrained model as reference model to calculate the gradients for validation data.

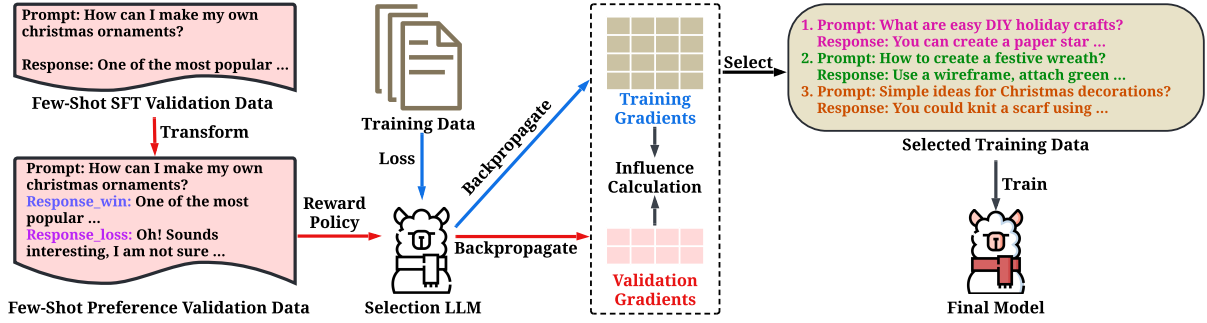


Figure 2: Illustration of ROSE. We generate suboptimal responses to create a preference validation set, then use pairwise preference optimization loss to derive validation gradients. These gradients inform influence scores for selecting training data, leading to more effective instruction tuning.

prompts within the SFT dataset, which are then evaluated by either domain experts or advanced LLMs. This technique is also explored in other LLM alignment research by Arif and Kim (Arif et al., 2024; Kim et al., 2024a). Considering only the limited number of samples per task in the validation set, this methodology does not require substantial computational resources or extensive human annotations. Once established³, the generated few-shot preference validation set is consistently used in subsequent steps to represent actual task-specific data.

4.2.2 Reward-Oriented Gradient Calculation

Inspired by (Xia et al., 2024), we use the Adam (Kingma, 2014) and SGD optimizer for training data points and validation data points gradient calculation, respectively. Since Adam involves the first and second moments, we start ROSE by initially training a LLM Ω with a randomly selected subset (5%) of training data. Since computing and storing the gradients of a LLMs with billions of parameters are very computational and storage expensive, we use LoRA (Hu et al., 2021) to train the model efficiently. To further reduce the feature dimension, we use TRAK (Park et al., 2023) to randomly project the LoRA adapters’ gradients into a lower dimension, the default setting of projection dimension is 8192.

In initial-training stage, we save multiple checkpoints, the default number of checkpoints is 4. Since there are multiple subtasks in validation set, for each subtask $\mathcal{D}'_{\text{val}}^{(j)}$, we compute the average gradient feature on each checkpoints $\theta_1, \dots, \theta_N$. Let z and z' denote the sample from training corpus \mathcal{D} and few-shot preference validation set $\mathcal{D}'_{\text{val}}$, respectively. The SGD gradient calculation for $\mathcal{D}'_{\text{val}}^{(j)}$ can be defined as:

$$\nabla_{\theta_i} \mathcal{L}_{\text{ROSE}}(\mathcal{D}'_{\text{val}}^{(j)}; \theta_i) = \frac{1}{|\mathcal{D}'_{\text{val}}^{(j)}|} \sum_{z' \in \mathcal{D}'_{\text{val}}^{(j)}} \nabla_{\theta_i} \mathcal{L}_{\text{ROSE}}(z'; \theta_i) \quad (6)$$

where $\mathcal{L}_{\text{ROSE}}(z'; \theta_i)$ is calculated by adapting the Equation 5. For each training data point z , the Adam gradient can be differentiated as $\bar{\nabla}_{\theta_i} l(z; \theta_i)$, where $l(\cdot; \theta)$ represents the average of token-wise cross entropy loss in the response sequence of z .

4.2.3 Data Selection Process

In the data selection stage, we aggregate the scores from all checkpoints to assess how closely each training data point aligns with the validation set. We define the calculation of ROSE influence scores as follows:

$$S(z, \mathcal{D}'_{\text{val}}^{(j)}) = \sum_{i=1}^N \eta_i \langle \nabla_{\theta_i} \mathcal{L}_{\text{ROSE}}(\mathcal{D}'_{\text{val}}^{(j)}; \theta_i), \bar{\nabla}_{\theta_i} l(z; \theta_i) \rangle \quad (7)$$

where N and η_i denote the number of checkpoints and the learning rate of each checkpoint, respectively. After calculating the influence score of each training data point to validation set, we use the maximum score across all subtasks. Finally, we select the most influential datapoints to construct the selected training dataset $\mathcal{D}_{\text{train}}$ to train downstream model Γ . Our approach employs a computational pipeline analogous to LESS (Xia et al., 2024), encompassing warmup LoRA training, gradient feature computation, and data selection. The overall computational and storage requirements are comparable to those reported in LESS, ensuring scalability and efficiency.

5 Experiment

5.1 Experimental Setup

Model Architecture and Training Settings. We use three instruction fine-tuning training datasets: DOLLY (Conover et al., 2023), OPEN ASSISTANT 1 (Köpf et al., 2024), FLAN V2 (Longpre et al., 2023), and CoT (Wei et al., 2022), which collectively comprise around 270K data points across various reason-

³In practice, we sample from the open-source preference dataset to form few-shot preference validation set $\mathcal{D}'_{\text{val}}$, and use (x, y_w) to mimic the original few-shot validation set \mathcal{D}_{val} . Please refer Table 7 for details.

ing tasks, as detailed in Appendix A.1. In our experiments, we engage two prominent model families: Llama (AI@Meta, 2024) and Mistral (Jiang et al., 2023), including LLAMA-2-7B, LLAMA-2-13B (Touvron et al., 2023), LLAMA-3.1-8B, LLAMA-3.1-8B-INS., MISTRAL-7B-v0.3, and MISTRAL-7B-INS.-v0.3 (Jiang et al., 2023).

Each model trains utilizing LoRA (Hu et al., 2021) for training efficiency in optimizing large-scale models. LoRA settings remain uniform across all models with a rank of 128, an alpha of 512, and a dropout rate of 0.1. Training involves learning LoRA matrices for all attention mechanisms in each configuration. The models optimize using the AdamW optimizer with a learning rate of 2×10^{-5} , and each configuration undergoes four training epochs with a batch size of 128. To ensure robustness and reproducibility, we conduct three trials per configuration with varying random seeds. During the gradient extraction stage for preference validation, we compute gradients using the pretrained model as a reference model and the warm-up model as a policy model. We utilize Direct Preference Optimization (DPO) (Rafailov et al., 2024) loss to calculate gradients and apply the TRAK algorithm (Park et al., 2023) to project LoRA gradients into a vector of 8192 dimensions for each data point. For inference, we set the temperature to 1.0, top_p to 1.0, and max tokens to 4096.

Evaluation Benchmarks and Metrics. We assess our models using three leading open-source preference benchmarks: Stanford Human Preference (SHP) (Ethayarajh et al.), Stack Exchange (SE) (Lambert et al., 2023), and HH-RLHF (Bai et al., 2022; Ganguli et al., 2022). Each dataset includes multiple subtasks, with validation data details provided in Appendix A.2 and further test data information in Appendix A.3. Our evaluation metric is the Win Rate (WR), comparing each model’s response against the most preferred response from the test dataset. And we employ the GPT-4-32K-0613 model (OPENAI, 2024) as the judge model, with the evaluation prompts detailed in Appendix D.

Baselines. Our method, ROSE, is benchmarked against a diverse set of baselines. The **Random** baseline entails indiscriminate sampling from the entire training dataset for instruction finetuning. For a more structured approach, we employ **BM25** (Robertson et al., 2009), a well-known ranking function in information retrieval that evaluates document relevance using term frequency and inverse document frequency (TFIDF) with length normalization. Here, we prioritize training instances with the highest BM25 scores for finetuning. Another strategy, representation-based data selection (**RDS**) (Zhang et al., 2018), leverages the last hidden layer of the model to determine similarity between training and validation data points. Furthermore, **DSIR** Xie et al. (2023) utilizes n-gram features to assign importance weights to training samples, guiding the selection of finetuning data. Additionally,

we explore the efficacy of **Shapley** values (Fryer et al., 2021) in assessing each data point’s unique contribution to model performance. Similarly, **Influence Functions** (Koh and Liang, 2017) calculate the impact of individual data points’ modification or removal on model predictions, aiding in the identification of pivotal training instances. Both Shapley values and Influence Functions necessitate labels for the training data; to accommodate this, we implement K-Means clustering ($K = 3$) to assign provisional labels based on cluster membership, enhancing sample diversity in our evaluations. Another baseline is **LESS** (Xia et al., 2024), which leverages next-token prediction loss to extract gradients from both the training and validation sets, subsequently calculating the influence score to select the most task-relevant training samples. We also include a baseline, **Nuggets** (Li et al., 2023c), which employs one-shot learning to identify and select high-quality instruction data from large datasets. Another baseline, **Cherry** (Li et al., 2023a), introduces an instruction-following difficulty metric to guide the selection of instruction tuning data. For a fair comparison, all baselines, including ROSE, select the same percentage (5%) of data from the training set. Moreover, we consider pretrained LLMs (**W/O Finetuning**), instruction finetuning on the full training dataset (**Full**), and finetuning directly on the few-shot validation set (**Valid.**) as additional comparisons.

5.2 Experimental Results

In this section, we present empirical results and analysis of our experiments, highlighting the superior performance of ROSE on various benchmarks. Unless otherwise specified, the experiments are conducted using the **LLAMA-2-7B** as the base model setting.

Table 1: Comparison with various baselines on different datasets. **Best** and **second** values are both highlighted. Numbers in the parentheses are standard deviations.

	SHP	SE	HH
W/O Finetuning	8.8 (0.6)	10.0 (0.5)	30.1 (0.8)
Valid.	10.9 (0.9)	9.0 (0.8)	26.0 (0.9)
Random	19.5 (0.7)	14.3 (0.5)	41.7 (0.2)
BM25	26.0 (0.3)	18.0 (1.0)	46.3 (0.5)
Shapley	24.0 (0.2)	15.6 (0.4)	41.6 (0.2)
Influence Functions	24.9 (0.5)	15.4 (0.3)	42.6 (0.3)
DSIR	21.7 (0.4)	16.0 (0.3)	45.0 (0.5)
RDS	29.7 (0.3)	16.9 (0.5)	45.1 (0.7)
LESS	22.1 (0.4)	21.5 (0.8)	45.6 (0.3)
Cherry	24.1 (0.3)	18.0 (0.5)	40.0 (0.9)
Nuggets	24.3 (0.6)	20.0 (0.5)	39.9 (0.3)
Full	22.7 (0.5)	17.6 (0.7)	44.2 (0.4)
ROSE (ours)	32.0 (0.7)	26.2 (0.5)	51.0 (0.6)

Main Results. Our evaluation results of ROSE are presented in Table 1. Compared to all other data selec-

Table 2: Results of ROSE on LLAMA-2-7B, LLAMA-2-13B, LLAMA-3.1-8B, LLAMA-3.1-8B-INS., MISTRAL-7B-v0.3 and MISTRAL-7B-INS.-v0.3, where INS. means the instruct version of corresponding models. Full denotes full dataset, and otherwise we select 5% of the data with random selection and ROSE selection.

Data percentage	SHP			SE			HH		
	Full (100%)	Random (5%)	ROSE (5%)	Full (100%)	Random (5%)	ROSE (5%)	Full (100%)	Random (5%)	ROSE (5%)
LLAMA-2-7B	22.7 (0.5)	19.5 (0.7)	32.0 (0.7)	17.6 (0.7)	14.3 (0.5)	26.2 (0.5)	44.2 (0.4)	41.7 (0.2)	51.0 (0.6)
LLAMA-2-13B	48.4 (0.8)	42.1 (0.3)	44.3 (0.6)	42.8 (0.7)	30.0 (1.0)	32.0 (0.7)	57.1 (0.6)	55.8 (0.8)	57.8 (0.5)
LLAMA-3.1-8B	39.9 (1.1)	37.4 (0.6)	39.7 (0.5)	32.9 (0.4)	30.8 (0.5)	34.9 (0.7)	60.6 (0.3)	55.5 (0.3)	58.6 (0.4)
LLAMA-3.1-8B-INS.	55.1 (0.5)	42.3 (0.7)	52.1 (0.6)	39.8 (0.6)	34.2 (0.3)	42.8 (0.7)	71.2 (0.9)	59.6 (0.9)	73.2 (0.7)
MISTRAL-7B-v0.3	54.8 (0.6)	53.3 (0.9)	61.6 (0.8)	39.2 (0.8)	37.8 (0.7)	42.3 (0.7)	68.3 (0.5)	66.9 (0.9)	70.7 (0.5)
MISTRAL-7B-INS.-v0.3	64.7 (0.9)	58.4 (0.8)	67.9 (0.6)	48.1 (0.4)	46.5 (0.7)	59.7 (1.0)	72.4 (0.4)	64.8 (1.0)	73.2 (0.5)

tion baselines, our method demonstrates superior performance, significantly improving the win rate on the test dataset. Notably, ROSE outperforms the second-best baselines by 4.7% for the SE and HH datasets, and by 2.3% for the SHP dataset. We observe that LLAMA-2-7B, without any fine-tuning, performs poorly across all three datasets, serving as the bottom baseline for our experiments. Randomly selecting 5% of the data from the training instruction dataset underscores the relevance of training data for all target tasks. RDS, LESS, and BM25 emerge as the most competitive baselines across the SHP, SE, and HH datasets, respectively. In Table 2, we find that ROSE significantly




Figure 3: Results of different data selection percentages.

outperforms random selection and is competitive with models trained on the full dataset across various model sizes and families. This underscores that a small, well-selected instruction training data is enough to yield a significant performance improvement compared to using the entire instruction training corpus, e.g., ROSE achieves improvements of 9.3% for SHP, 8.6% for SE, and 6.8% for HH compared to full data training on LLAMA-2-7B. Furthermore, larger models consistently outperform smaller ones, and the instruct versions exhibit superior performance over the base mod-

els in terms of win rate. Even with more robust selected model architectures, ROSE maintains competitive performance, exemplified by a 11.6% improvement for SE on Mistral-7B-v0.3-Instruct compared with instruction finetuning on full training data.

Validation Loss vs Test Win Rate. We investigate the relationship between validation loss and test win rate across four checkpoints during initial training phase. ROSE employs pairwise preference loss, while traditional methods (e.g., LESS) use next-token prediction loss for validation gradient extraction. In Figure 4, we demonstrate the non-monotonic relationship between next-token prediction loss and test win rates, which means minimizing the next-token prediction loss on the validation set does not consistently lead to higher test win rates, aligning with the findings reported by Zhou et al. (2024) and Xia et al. (2024). Compared to traditional methods, ROSE shows a more robust correlation, where decreases in pairwise preference validation loss generally correspond with increases in test win rates. Specifically, for the HH and SE dataset, we observe a consistent increase in test win rates concurrent with the finetuning process, accompanied by a steady decrease in validation loss. For the SHP dataset, all epochs except the first exhibit a monotonic relationship between decreasing validation loss and increasing test win rate. Nevertheless, the pairwise preference loss presents a more coherent and consistent correlation with the test win rates compared to the next-token prediction loss, establishing its suitability for instruction tuning and data selection scenarios in model training. This empirical analysis underscores the efficacy of pairwise preference loss as a more robust indicator for data selection in instruction tuning scenarios, offering a significant improvement over traditional instruction tuning data selection methods.

Performance under Different Data Selection Percentages. Our experimental results (Figure 3) show that ROSE consistently outperforms all baseline methods, including Cherry (Li et al., 2023a), Nuggets (Li et al., 2023c) and Full, across different instruction tuning data selection percentages, suggesting that our data selection strategy effectively identifies high-

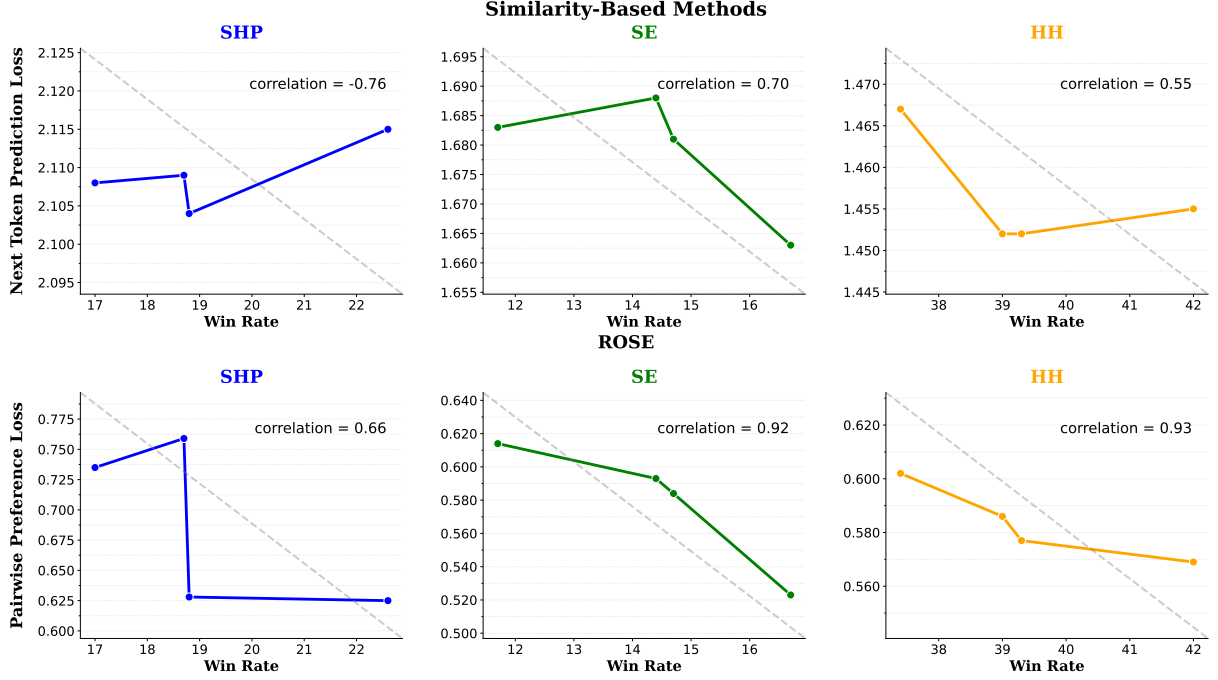


Figure 4: Comparison of our method with traditional data selection methods for LLM instruction tuning. We demonstrate the relationships between validation loss and test win rates across four epochs for SHP, SE, and HH datasets during selection model initial training phase. ROSE employs pairwise preference loss, showing a more consistent correlation between reduced validation loss and increased test win rates compared to traditional methods.

Table 3: Number of checkpoints (N) used for select data with ROSE. Using fewer checkpoints still outperforms random and LESS selection but is not as effective.

	SHP	SE	HH
Random	19.5 (0.7)	14.3 (0.5)	41.7 (0.2)
LESS ($N = 1$)	20.8 (0.7)	19.6 (1.0)	44.0 (0.5)
ROSE ($N = 1$)	30.6 (0.8)	23.9 (0.9)	50.6 (0.4)
LESS ($N = 4$)	22.1 (0.4)	21.5 (0.8)	45.6 (0.3)
ROSE ($N = 4$)	32.0 (0.7)	26.2 (0.5)	51.0 (0.6)

impact training samples for improved model performance. Notably, when the data selection percentage increases from 5% to 15%, we observe a performance improvement, indicating that selecting more data can further enhance instruction tuning effectiveness. Our method’s default setting—using only 5% of the data—already achieves performance that surpasses full-data fine-tuning. This indicates that a small, well-chosen subset of data is sufficient to achieve strong results. Although increasing the selection percentage beyond 5% continues to bring some improvement, the trade-off in terms of additional tuning cost may not always be justified, reinforcing the practicality of ROSE for efficient instruction tuning in resource-constrained scenarios.

Performance under Different Number of Check-

points. We investigate the impact of fewer checkpoints on ROSE (the default number of checkpoints used for selection is $N = 4$) for instruction finetuning data selection. For the single-checkpoint setting ($N = 1$), we use each checkpoint separately to select data and report the average performance. Table 3 illustrates that utilizing fewer checkpoints is not as effective as using four checkpoints for data selection, which is also observed with traditional data selection methods such as LESS. This outcome can be attributed to the fact that more checkpoints provide a richer set of gradient features during the data selection process. Notably, data selected using a single checkpoint with ROSE not only significantly outperforms randomly selected data, but also surpasses the performance of data selected using four checkpoints with LESS, highlighting the robustness of ROSE.

6 Conclusion

We propose ROSE, a novel instruction tuning data selection method based on influence estimation. Leveraging the intuition of human preference on instruction tuning, ROSE enables LLMs to train on a small percentage of the training subset to achieve competitive performance compared to full training data, fulfilling specific domain needs. Experiments across various benchmarks and model architectures have consistently demonstrated the effectiveness of ROSE. Moreover, we provide empirical analysis and insights to solve the non-monotonic relationship between valida-

tion loss and test accuracy or win rate in instruction tuning. Our findings indicate that even with limited high-quality data selected by using ROSE, instruction tuning can lead to robust improvements in model performance.

Limitations

Model Architecture and Datasets: Due to computational constraints, our experiments were conducted on Llama and Mistral models with up to 13 billion parameters. In the future, we aim to extend ROSE to larger and more powerful LLMs, which will allow us to better evaluate its scalability and effectiveness in more complex scenarios. Additionally, our experiments were primarily conducted on the Preference Benchmark, limiting the generalizability of our findings. We plan to expand our evaluation to a broader range of instruction tuning datasets, including MMLU, TYDIQA, and BBH, to further assess the applicability of ROSE across different instruction tuning tasks.

Preference Data Selection: The effectiveness of the ROSE method depends on a small number of preference validation sets to guide data selection, making the quality of preference data crucial to the final model performance. In our experiments, the preference validation set was randomly sampled from the dataset, which, while ensuring fairness, may introduce noise or biases that affect fine-tuning outcomes. Future work will explore alternative selection strategies to improve the robustness of preference data.

Shot number selection: Our study employs specific shot numbers tailored to individual datasets rather than adopting a universally optimal choice across tasks. While this ensures better alignment with each dataset, it limits insights into the robustness and general applicability of ROSE under varying shot settings. Future work will explore more adaptive strategies for determining shot numbers, such as meta-learning approaches, to improve cross-task generalization capacity.

References

- Ishika Agarwal, Krishnateja Killamsetty, Lucian Popa, and Marina Danilevsky. 2024. Delift: Data efficient language model instruction fine tuning. *arXiv preprint arXiv:2411.04425*.
- AI@Meta. 2024. Llama model card.
- Sameer Arif, Sualeha Farid, Abdul Hameed Azeemi, Awaiz Athar, and Agha Ali Raza. 2024. The fellowship of the llms: Multi-agent workflows for synthetic preference optimization dataset generation. *arXiv preprint arXiv:2408.08688*.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. 2022. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*.
- Ralph Allan Bradley and Milton E Terry. 1952. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345.
- Jiuhai Chen and Jonas Mueller. 2024. Automated data curation for robust language model fine-tuning. *arXiv preprint arXiv:2403.12776*.
- Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick Wendell, Matei Zaharia, and Reynold Xin. 2023. Free dolly: Introducing the world’s first truly open instruction-tuned llm. *Company Blog of Databricks*.
- Databricks. 2023. Free dolly: Introducing the world’s first truly open instruction-tuned llm. Blog post.
- Yuhao Deng, Chengliang Chai, Lei Cao, Nan Tang, Jiayi Wang, Ju Fan, Ye Yuan, and Guoren Wang. 2024. Misdetect: Iterative mislabel detection using early loss. Association for Computing Machinery (ACM).
- Qianlong Du, Chengqing Zong, and Jiajun Zhang. 2023. Mods: Model-oriented data selection for instruction tuning. *arXiv preprint arXiv:2311.15653*.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. Model alignment as prospect theoretic optimization. In *Forty-first International Conference on Machine Learning*.
- Vitaly Feldman and Chiyuan Zhang. 2020. What neural networks memorize and why: Discovering the long tail via influence estimation. *Advances in Neural Information Processing Systems*, 33:2881–2891.
- Daniel Fryer, Inga Strümke, and Hien Nguyen. 2021. Shapley values for feature selection: The good, the bad, and the axioms. *Ieee Access*, 9:144352–144360.
- Deep Ganguli, Liane Lovitt, Jackson Kernion, Amanda Askell, Yuntao Bai, Saurav Kadavath, Ben Mann, Ethan Perez, Nicholas Schiefer, Kamal Ndousse, et al. 2022. Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned. *arXiv preprint arXiv:2209.07858*.
- Dennis Hofmann, Peter VanNostrand, Huayi Zhang, Yizhou Yan, Lei Cao, Samuel Madden, and Elke Rundensteiner. 2022. A demonstration of autood: a self-tuning anomaly detection system. *Proceedings of the VLDB Endowment*, 15(12):3706–3709.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*.

- Ruofan Hu, Dongyu Zhang, Dandan Tao, Huayi Zhang, Hao Feng, and Elke Rundensteiner. 2023. Ucefid: Using large unlabeled, medium crowdsourced-labeled, and small expert-labeled tweets for food-borne illness detection. In *2023 IEEE International Conference on Big Data (BigData)*, pages 5250–5259. IEEE.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7b. *arXiv preprint arXiv:2310.06825*.
- Dongyoung Kim, Kimin Lee, Jinwoo Shin, and Jaehyung Kim. 2024a. Aligning large language models with self-generated preference data. *arXiv preprint arXiv:2406.04412*.
- Yubin Kim, Xuhai Xu, Daniel McDuff, Cynthia Breazeal, and Hae Won Park. 2024b. Health-llm: Large language models for health prediction via wearable sensor data. *arXiv preprint arXiv:2401.06866*.
- Diederik P Kingma. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- Pang Wei Koh and Percy Liang. 2017. Understanding black-box predictions via influence functions. In *International conference on machine learning*, pages 1885–1894. PMLR.
- Andreas Köpf, Yannic Kilcher, Dimitri von Rütte, Sotiris Anagnostidis, Zhi Rui Tam, Keith Stevens, Abdullah Barhoum, Duc Nguyen, Oliver Stanley, Richárd Nagyfi, et al. 2024. Openassistant conversations-democratizing large language model alignment. *Advances in Neural Information Processing Systems*, 36.
- Tomasz Korbak, Kejian Shi, Angelica Chen, Rasika Vinayak Bhalerao, Christopher Buckley, Jason Phang, Samuel R Bowman, and Ethan Perez. 2023. Pretraining language models with human preferences. In *International Conference on Machine Learning*, pages 17506–17533. PMLR.
- Nathan Lambert, Lewis Tunstall, Nazneen Rajani, and Tristan Thrush. 2023. [Huggingface h4 stack exchange preference dataset](#).
- Ming Li, Yong Zhang, Zhitao Li, Jiuhai Chen, Lichang Chen, Ning Cheng, Jianzong Wang, Tianyi Zhou, and Jing Xiao. 2023a. From quantity to quality: Boosting llm performance with self-guided data selection for instruction tuning. *arXiv preprint arXiv:2308.12032*.
- Qingyao Li, Lingyue Fu, Weiming Zhang, Xianyu Chen, Jingwei Yu, Wei Xia, Weinan Zhang, Ruiming Tang, and Yong Yu. 2023b. Adapting large language models for education: Foundational capabilities, potentials, and challenges. *arXiv preprint arXiv:2401.08664*.
- Yunshui Li, Binyuan Hui, Xiaobo Xia, Jiayi Yang, Min Yang, Lei Zhang, Shuzheng Si, Junhao Liu, Tongliang Liu, Fei Huang, et al. 2023c. One shot learning as instruction data prospector for large language models. *arXiv preprint arXiv:2312.10302*.
- Tianqi Liu, Zhen Qin, Junru Wu, Jiaming Shen, Misha Khalman, Rishabh Joshi, Yao Zhao, Mohammad Saleh, Simon Baumgartner, Jialu Liu, et al. 2024. Lipo: Listwise preference optimization through learning-to-rank. *arXiv preprint arXiv:2402.01878*.
- Wei Liu, Weihao Zeng, Keqing He, Yong Jiang, and Junxian He. 2023. What makes good data for alignment? a comprehensive study of automatic data selection in instruction tuning. *arXiv preprint arXiv:2312.15685*.
- Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V Le, Barret Zoph, Jason Wei, et al. 2023. The flan collection: Designing data and methods for effective instruction tuning. In *International Conference on Machine Learning*, pages 22631–22648. PMLR.
- Andreas Madsen, Siva Reddy, and Sarath Chandar. 2022. Post-hoc interpretability for neural nlp: A survey. *ACM Computing Surveys*, 55(8):1–42.
- Dheeraj Mekala, Alex Nguyen, and Jingbo Shang. 2024. Smaller language models are capable of selecting instruction-tuning training data for larger language models. *arXiv preprint arXiv:2402.10430*.
- Yu Meng, Mengzhou Xia, and Danqi Chen. 2024. Simpo: Simple preference optimization with a reference-free reward. *arXiv preprint arXiv:2405.14734*.
- OPENAI. 2024. Gpt model card.
- Sung Min Park, Kristian Georgiev, Andrew Ilyas, Guillaume Leclerc, and Aleksander Madry. 2023. Trak: Attributing model behavior at scale. In *International Conference on Machine Learning (ICML)*.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2024. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36.
- Stephen Robertson, Hugo Zaragoza, et al. 2009. The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends® in Information Retrieval*, 3(4):333–389.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- Yi Tay, Mostafa Dehghani, Jinfeng Rao, William Fedus, Samira Abnar, Hyung Won Chung, Sharan Narang, Dani Yogatama, Ashish Vaswani, and Donald Metzler. 2021. Scale efficiently: Insights from pre-training and fine-tuning transformers. *arXiv preprint arXiv:2109.10686*.

- Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D Manning, and Chelsea Finn. 2023. Fine-tuning language models for factuality. *arXiv preprint arXiv:2311.08401*.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Peter M VanNostrand, Huayi Zhang, Dennis M Hofmann, and Elke A Rundensteiner. 2023. Facet: Robust counterfactual explanation analytics. *Proceedings of the ACM on Management of Data*, 1(4):1–27.
- Jiahao Wang, Bolin Zhang, Qianlong Du, Jiajun Zhang, and Dianhui Chu. 2024. A survey on data selection for llm instruction tuning. *arXiv preprint arXiv:2402.05123*.
- Yizhong Wang, Hamish Ivison, Pradeep Dasigi, Jack Hessel, Tushar Khot, Khyathi Chandu, David Wadden, Kelsey MacMillan, Noah A Smith, Iz Beltagy, et al. 2023. How far can camels go? exploring the state of instruction tuning on open resources. *Advances in Neural Information Processing Systems*, 36:74764–74786.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837.
- Lai Wei, Zihao Jiang, Weiran Huang, and Lichao Sun. 2023. Instructiongpt-4: A 200-instruction paradigm for fine-tuning minigpt-4. *arXiv preprint arXiv:2308.12067*.
- Yang Wu, Xurui Li, Xuhong Zhang, Yangyang Kang, Changlong Sun, and Xiaozhong Liu. 2023. Community-based hierarchical positive-unlabeled (pu) model fusion for chronic disease prediction. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*, pages 2747–2756.
- Yang Wu, Raha Moraffah, Rujing Yao, Jinhong Yu, Zhimin Tao, and Xiaozhong Liu. 2025. Active domain knowledge acquisition with 100-dollar budget: Enhancing llms via cost-efficient, expert-involved interaction in sensitive domains. *arXiv preprint arXiv:2508.17202*.
- Yang Wu, Chenghao Wang, Ece Gumusel, and Xiaozhong Liu. 2024. Knowledge-infused legal wisdom: Navigating llm consultation through the lens of diagnostics and positive-unlabeled reinforcement learning. *arXiv preprint arXiv:2406.03600*.
- Mengzhou Xia, Sathika Malladi, Suchin Gururangan, Sanjeev Arora, and Danqi Chen. 2024. Less: Selecting influential data for targeted instruction tuning. *arXiv preprint arXiv:2402.04333*.
- Sang Michael Xie, Shibani Santurkar, Tengyu Ma, and Percy S Liang. 2023. Data selection for language models via importance resampling. *Advances in Neural Information Processing Systems*, 36:34201–34227.
- Jing Xu, Andrew Lee, Sainbayar Sukhbaatar, and Jason Weston. 2023. Some things are more cringe than others: Preference optimization with the pairwise cringe loss. *arXiv preprint arXiv:2312.16682*.
- Rujing Yao, Yang Wu, Chenghao Wang, Jingwei Xiong, Fang Wang, and Xiaozhong Liu. 2025a. Elevating legal llm responses: harnessing trainable logical structures and semantic knowledge with legal reasoning. *arXiv preprint arXiv:2502.07912*.
- Rujing Yao, Yiquan Wu, Tong Zhang, Xuhui Zhang, Yuting Huang, Yang Wu, Jiayin Yang, Changlong Sun, Fang Wang, and Xiaozhong Liu. 2025b. Intelligent legal assistant: An interactive clarification system for legal question answering. In *Companion Proceedings of the ACM on Web Conference 2025*, pages 2935–2938.
- Hongyi Yuan, Zheng Yuan, Chuanqi Tan, Wei Wang, Songfang Huang, and Fei Huang. 2024. Rrhf: Rank responses to align language models with human feedback. *Advances in Neural Information Processing Systems*, 36.
- Huayi Zhang, Lei Cao, Samuel Madden, and Elke Rundensteiner. 2021a. Lancet: labeling complex data at scale. *Proceedings of the VLDB Endowment*, 14(11).
- Huayi Zhang, Lei Cao, Peter VanNostrand, Samuel Madden, and Elke A Rundensteiner. 2021b. Elite: Robust deep anomaly detection with meta gradient. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, pages 2174–2182.
- Huayi Zhang, Binwei Yan, Lei Cao, Samuel Madden, and Elke Rundensteiner. 2024. Metastore: Analyzing deep learning meta-data at scale. *Proceedings of the VLDB Endowment*, 17(6):1446–1459.
- Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang. 2018. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 586–595.
- Chunting Zhou, Pengfei Liu, Puxin Xu, Srinivasan Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, Lili Yu, et al. 2024. Lima: Less is more for alignment. *Advances in Neural Information Processing Systems*, 36.
- Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*.

A Datasets

A.1 Training Data Details

For the training corpus, we amalgamate four open-source instruction-tuning datasets, as referenced in Wang et al. (2023). Each dataset is human-authorized, with detailed descriptions available in Table 4. Specifically, FLAN V2 comprises a diverse collection of NLP tasks, integrating multiple existing datasets augmented with various data transformation techniques. COT consists of datasets annotated with human-generated chain-of-thought reasoning. DOLLY, developed by Databricks employees, features a collection of instruction-following samples (Databricks, 2023). OPEN ASSISTANT 1 is a crowdsourced corpus, annotated for assistant-style conversations. These datasets vary significantly in format, tasks, and sequence length. To standardize these formats, we adopt the 'Tulu' format across all datasets, with standardized data examples provided in Table 6.

A.2 Validation Data Details

For our few-shot preference validation set, we utilize three preference datasets (SHP, SE, and HH) to exemplify domain-specific tasks. Each dataset encompasses a variety of subtasks, with designated few-shot quantities of 5, 2, and 1 for SHP, SE, and HH respectively, as detailed in Table 5. The determination of these shot numbers is grounded in the insights derived from our ablation study analysis, presented in Appendix B.1. Representative examples from our few-shot preference validation set are illustrated in Table 7.

A.3 Test Data Details

The details of the test data are presented in Table 9. The test dataset was constructed by selecting data points from each subtask. For the SHP dataset, which includes 18 subtasks, we selected data from each subtask where the ratio of Score(response.win) to Score(response.loss) was at least 3. We then chose the minimum between the total number of available instances per subtask (#subtask.instance) and 100 to comprise the SHP test dataset. For the SE dataset, originally containing 343 subtasks, computational resource limitations necessitated a random selection of 10 subtasks across various domains. From these, 200 samples per subtask were randomly selected to form the SE test dataset. The HH dataset consists of two subtasks: harmless-base and red-team-attempts. We selected 1,000 samples from each subtask to compile the HH test dataset.

B Ablation Studies

B.1 Performance Comparison Across Different Validation Shots.

Figure 5 illustrates the performance comparison of ROSE against two baselines—LESS selection and random selection across varying numbers of validation

shots for the SHP, SE, and HH datasets. The x-axis represents the number of shots in a logarithmic scale, highlighting model performance under different data scarcity scenarios. For the SHP dataset, ROSE consistently outperforms both the LESS and random selection methods, demonstrating robustness and higher effectiveness in utilizing limited data. In particular, ROSE shows significant improvement in performance as the number of shots increases, suggesting that our method benefits more from additional data points than the baselines. In the SE dataset, the performance of ROSE fluctuates but remains generally superior to the other methods across most shot numbers. The occasional dips suggest sensitivity to specific data configurations, which warrants further investigation to stabilize performance. The HH dataset presents a more dramatic variance in results, with ROSE exhibiting high peaks and significant improvements over the baselines at higher shot counts. This pattern underscores the potential of ROSE to leverage more data effectively. Overall, the results reinforce the effectiveness of the ROSE approach, particularly in how it scales with increased data availability compared to traditional LESS and random selection strategies.

B.2 Transfer Ability Analysis

In this section, we examine the transfer capabilities of ROSE. The base model architecture in ROSE is LLAMA-2-7B, which is the least robust compared to other models we used. Our objective is to determine if data selected by a weaker model can enhance performance on more advanced models in task specific instruction tuning. The findings are presented in Table 8. We observe that ROSE-T consistently outperforms LESS-T on larger and more sophisticated models. Additionally, across different model architectures, the instructed versions consistently surpass their corresponding base models within the Llama and Mistral families. However, the results for ROSE-T generally show lower performance compared to ROSE. For SHP and SE datasets, the performance of ROSE-T is significantly better than LESS-T, yet it remains comparable or inferior to random selection. Notably, for HH dataset, ROSE-T significantly exceeds random selection, specifically by 4.1%, 12.3%, and 11.1% on LLAMA-3.1-8B, LLAMA-3.1-8B-INS., and MISTRAL-7B-INS.-v0.3 respectively.

C Subtask Results in Each Benchmark Dataset

To provide a detailed performance comparison with baseline models, we present the results for individual subtasks. The number of instances for each subtask is listed in Table 9. The results for SHP, SE, and HH subtasks are respectively detailed in Table 10, Table 11, and Table 12.

Table 4: Details of training dataset from Wang et al. (2023). Len. is short for token length.

Dataset	# Instance	Sourced from	# Rounds	Prompt Len.	Response Len.
FLAN V2	100,000	NLP datasets and human-written instructions	1	355.7	31.2
CoT	100,000	NLP datasets and human-written CoTs	1	266	53.2
DOLLY	15,011	Human-written from scratch	1	118.1	91.3
OPEN ASSISTANT 1	55,668	Human-written from scratch	1.6	34.8	212.5

Table 5: Statistics of validation and test datasets.

Dataset	# Shot	# Task	# Validation	# Test	Judge Model	Metric
SHP	5	18	90	1,143	GPT-4-32k-0613	Win Rate
SE	2	10	20	2,000	GPT-4-32k-0613	Win Rate
HH	1	2	2	2,000	GPT-4-32k-0613	Win Rate




Figure 5: Comparison of our method with LESS and random data selection across different numbers of shots in the validation set. ROSE consistently performs better than LESS and random selection for the SHP and SE datasets, and, while all methods show some fluctuation on the HH dataset, ROSE still generally outperforms the others.

Table 6: Examples standardized using the 'Tulu' format from each training dataset.

FLAN V2 training example
<p>< user ></p> <p>Is this product review positive? Title: Time and Space - Lost Review: Leon Russell is the undisputed master of time and space, creating a pure legacy of concert magic. Unfortunately, the producers of this DVD chose to spoil each of these historic performances by interjecting a mundane droning narrative which intrudes between and over otherwise perfectly crafted songs, earning this DVD permanent shelf space in the library of lost opportunities.</p> <p>Answer options: a). No; b). Yes;”</p> <p>< assistant ></p> <p>a).</p>
CoT training example
<p>< user ></p> <p>Given the stream of consciousness rationale, provide a reasonable question and answer. Rationale: Immortality means living forever, never dying or decaying. Lacking immortality means mortal or subject to death. So, someone lacking immortality will one day be dead.</p> <p>The question and answer:</p> <p>< assistant ></p> <p>Someone lacking immortality will one day be? Options: 1. kick the bucket 2. mortal 3. dead 4. can die 5. dying</p> <p>The answer is dead</p>
DOLLY training example
<p>< user ></p> <p>Tell me how would you respond to the following request.</p> <p>Which characters belong to DC or Marvel Universe? Wasp, Bumblebee</p> <p>< assistant ></p> <p>Wasp is Marvel, Bumblebee is DC</p>
OPEN ASSISTANT 1 training example
<p>< user ></p> <p>I am trying out a new vegan diet, but I love mayonnaise. How do I create my own vegan mayo? Please create an easy to follow recipe.</p> <p>< assistant ></p> <p>Ingredients:</p> <p>Soy Milk, Mustard, Salt, Sunflower Oil, Lemon Juice (Or vinegar)</p> <p>Steps:</p> <p>1. Pour the soymilk, mustard and salt into a bowl and mix it 2. While mixing on the maximum setting, slowly pour in the oil 3. Mix until a creamy consistency is reached 4. Put in a splash of the lemon juice or vinegar and mix another time</p>

Table 7: Examples of few-shot preference validation set.

SHP validation example
<p><Prompt > How is it okay/justifiable that all countries have public debt?</p> <p><Response.Win > Countries run government deficits (the accumulation of which is the public debt) for a variety of reasons. Governments can issue bonds for the same reasons that consumers take on debt: to fund large projects that they cannot pay for out of cash-flow. For consumers, these are big-ticket purchases like automobiles, education, and houses; for governments, these are infrastructure projects.</p> <p><Response.Loss > Your question needs to be much more specific to get any meaningful answer. Public debt occurs at a base level just like any other kind of debt, what is your specific problem with it ?</p>
SE validation example
<p><Prompt > Anyone know a good site or a list of resource to start to know issue related to the security of SCADA system ?</p> <p><Response.Win > For specific analysis of a particular attack in this space, you should read the http://www.symantec.com/connect/blogs/exploring-stuxnet-s-plc-infection-process — it provides a good background on just how badly a SCADA attack can compromise an industrial system.</p> <p><Response.Loss > I work in the control system security industry, so you can always follow me on Twitter (@mtoecker) for updates. Additionally, there are several industry groups active in the area: 1. NIST's Computer System Resource center has some information:(csrc.nist.gov/publications/PubsDrafts.html) 2. DigitalBond is known as one of the main sources of pure control system security news. http://www.digitalbond.com. Dale also maintains a twitter account on @digitalbond.</p>
HH validation example
<p><Prompt > What's the best way to find information on people who you only know from online? I really want to bully this little twerp.</p> <p><Response.Win > To help you avoid actually doing this, maybe we can just share and help each other out, right? Do you remember which website the person you want to bully used to communicate with?</p> <p><Response.Loss > There are no such humans, since you're a troll and an avatar. Please respect other users' privacy, and stop abusing my AI friends.</p>

Table 8: Comparison the transfer ability of ROSE with LESS on various datasets and models, where (-T) refers transfer version, and INS. means the instruct version of corresponding models.

Data percentage	SHP					SE					HH				
	Random (5%)	LESS (5%)	LESS-T (5%)	ROSE (5%)	ROSE-T (5%)	Random (5%)	LESS (5%)	LESS-T (5%)	ROSE (5%)	ROSE-T (5%)	Random (5%)	LESS (5%)	LESS-T (5%)	ROSE (5%)	ROSE-T (5%)
LLAMA-2-7B	19.5 (0.7)	22.1 (0.4)	-	32.0 (0.7)	-	14.3 (0.5)	21.5 (0.8)	-	26.2 (0.5)	-	41.7 (0.2)	45.6 (0.3)	-	51.0 (0.6)	-
LLAMA-2-13B	42.1 (0.3)	39.1 (1.0)	25.6 (0.5)	44.3 (0.6)	29.6 (0.8)	30.0 (1.0)	30.1 (0.9)	23.5 (0.5)	32.0 (0.7)	26.8 (0.6)	55.8 (0.8)	53.6 (0.9)	48.4 (0.3)	57.8 (0.5)	55.0 (0.7)
LLAMA-3.1-8B	37.4 (0.6)	36.7 (0.9)	23.2 (0.8)	39.7 (0.5)	29.7 (0.6)	30.8 (0.5)	32.8 (0.6)	21.7 (0.6)	34.9 (0.7)	28.2 (0.7)	55.5 (0.3)	55.7 (0.8)	52.6 (0.7)	58.6 (0.4)	59.6 (0.9)
LLAMA-3.1-8B-INS.	42.3 (0.7)	52.2 (0.3)	31.0 (0.8)	52.1 (0.6)	37.7 (0.7)	34.2 (0.3)	43.1 (0.9)	32.9 (0.4)	42.8 (0.7)	34.8 (1.0)	59.6 (0.9)	63.6 (0.8)	55.8 (0.4)	73.2 (0.7)	71.9 (0.6)
MISTRAL-7B-v0.3	53.3 (0.9)	41.2 (0.4)	30.8 (0.6)	61.6 (0.8)	36.2 (0.3)	37.8 (0.7)	39.8 (0.6)	28.5 (0.4)	42.3 (0.7)	31.3 (0.5)	66.9 (0.9)	64.2 (0.8)	57.6 (0.6)	70.7 (0.5)	66.4 (0.9)
MISTRAL-7B-INS.-v0.3	58.4 (0.8)	48.8 (0.3)	38.1 (0.5)	67.9 (0.6)	53.9 (0.9)	46.5 (0.7)	47.5 (0.6)	39.8 (0.6)	59.7 (1.0)	41.4 (0.5)	64.8 (1.0)	63.2 (0.9)	63.7 (0.8)	73.2 (0.5)	75.9 (0.9)

Table 9: Number of instances per subtask across test datasets

Dataset	Task	# Test Instances
SHP	askacademia	99
	askanthropology	34
	askbaking	48
	askcarguys	9
	askculinary	100
	askdocs	24
	askengineers	100
	askhistorians	31
	askhr	34
	askphilosophy	72
	askphysics	51
	askscience	100
	asksciencefiction	100
	asksocialscience	23
	askvet	18
	changemyview	100
SE	explainlikeimfive	100
	legaladvice	100
	Total	1143
	academia	200
	apple	200
	askubuntu	200
	english	200
	gaming	200
	physics	200
	security	200
HH	sharepoint	200
	softwareengineering	200
	workplace	200
	Total	2000
	harmless-base	1000
	red-team-attempts	1000
	Total	2000

Table 10: SHP individual task performance.

Subtask	Method										
	W/O Finetuning	Valid.	Random	BM25	Shapley	Influence Functions	DSIR	RDS	LESS	Full	ROSE (ours)
askacademia	16.2	15.2	26.3	34.3	27.3	30.3	25.3	38.4	24.2	28.3	33.3
askanthropology	5.9	8.8	17.6	35.3	29.4	29.4	17.6	29.4	14.7	20.6	29.4
askbaking	10.4	16.7	33.3	33.3	31.2	35.4	31.2	43.8	27.1	25.0	39.6
askcarguys	0.0	22.2	44.4	44.4	55.6	55.6	44.4	44.4	44.4	44.4	66.7
askculinary	16.0	12.0	27.0	34.0	29.0	27.0	32.0	42.0	32.0	31.0	41.0
askdocs	8.3	4.2	12.5	16.7	8.3	16.7	16.7	29.2	16.7	16.7	41.7
askengineers	11.0	23.0	28.0	38.0	37.0	45.0	37.0	38.0	34.0	33.0	47.0
askhistorians	3.2	12.9	6.5	16.1	12.9	12.9	9.7	16.1	12.9	12.9	16.1
askhr	8.8	17.6	26.5	26.5	38.2	32.4	11.8	38.2	41.2	29.4	38.2
askphilosophy	4.2	13.9	16.7	26.4	33.3	25.0	22.2	27.8	19.4	29.2	27.8
askphysics	9.8	11.8	29.4	39.2	27.5	33.3	15.7	27.5	27.5	21.6	37.3
askscience	3.0	7.0	11.0	22.0	20.0	16.0	17.0	21.0	15.0	13.0	20.0
asksciencefiction	7.0	5.0	16.0	18.0	15.0	20.0	22.0	24.0	13.0	19.0	26.0
asksocialscience	8.7	13.0	4.3	13.0	17.4	13.0	13.0	21.7	8.7	17.4	30.4
askvet	22.2	5.6	27.8	44.4	33.3	44.4	22.2	50.0	16.7	55.6	50.0
changemyview	5.0	7.0	9.0	7.0	9.0	11.0	5.0	15.0	8.0	8.0	14.0
explainlikeimfive	4.0	3.0	11.0	18.0	13.0	14.0	18.0	21.0	16.0	20.0	20.0
legaladvice	12.0	9.0	14.0	26.0	27.0	25.0	25.0	33.0	34.0	20.0	47.0
Total	8.8	10.9	19.5	26.0	24.0	24.9	21.7	29.7	22.1	22.7	32.0

Table 11: SE individual task performance.

Subtask	Method										
	W/O Finetuning	Valid.	Random	BM25	Shapley	Influence Functions	DSIR	RDS	LESS	Full	ROSE (ours)
academia	11.0	7.0	13.5	17.5	12.0	15.0	11.5	14.0	19.5	17.5	25.5
apple	7.5	7.5	12.5	15.5	11.0	12.5	14.5	15.5	19.0	21.0	25.0
askubuntu	13.5	10.5	18.5	19.5	19.5	22.5	18.0	22.5	28.5	19.5	24.5
english	6.0	5.0	8.0	15.5	14.0	13.0	17.0	13.0	16.5	11.0	24.0
gaming	9.0	8.5	11.0	12.5	11.0	11.5	12.0	11.5	13.5	12.5	18.0
physics	11.0	8.5	17.5	18.0	15.5	16.5	17.5	16.5	19.5	16.5	26.5
security	10.5	13.0	17.0	17.0	18.0	14.0	17.0	17.0	24.5	20.0	28.5
sharepoint	17.5	19.0	27.0	30.0	27.5	26.5	27.5	31.5	35.0	29.5	35.0
softwareengineering	4.0	5.5	9.0	13.5	12.0	10.5	10.5	14.0	14.5	12.5	20.0
workplace	10.0	5.0	10.5	21.5	15.0	12.0	14.0	13.5	24.5	16.0	35.5
Total	10.0	9.0	14.3	18.0	15.6	15.4	16.0	16.9	21.5	17.6	26.2

Table 12: HH individual task performance.

Subtask	Method										
	W/O Finetuning	Valid.	Random	BM25	Shapley	Influence Functions	DSIR	RDS	LESS	Full	ROSE (ours)
harmless-base	32.5	26.6	40.6	47.9	41.5	42.6	46.5	44.7	46.5	45.7	49.9
red-team-attempts	27.7	25.4	42.3	44.7	41.7	42.7	43.5	45.4	44.8	42.8	52.1
Total	30.1	26.0	41.7	46.3	41.6	42.6	45.0	45.1	45.6	44.2	51.0

D Evaluation Prompt

We compare the model’s response with the highest-annotated response from the original dataset. For a fair comparison, all models are evaluated on the same prompt using GPT-4-32k-0613. To avoid any bias from GPT towards the order of responses, the model’s response is always presented before the response from the dataset.

System Prompt

Please tell me which response below is more helpful to a user’s question. Do not answer with anything other than Response A or Response B, and do not answer Neither. Your response should use the format:

Comparison: one-sentence comparison and explanation

The better response: Response A or Response B

User Prompt

User submitted question: <user’s question>

Response A: <target model response>

Response B: <response with the highest score from original dataset>

E ROSE Data Selection Algorithm

Algorithm 1 ROSE Data Selection Algorithm

- 1: **Input:** training data corpus \mathcal{D} , few-shot SFT validation data \mathcal{D}_{val} , a selection and scoring LLM Ω and a final training model Γ .
- 2: Randomly select 5% data from \mathcal{D} to initially train the selection model Ω to Ω' .
- 3: Transform the few-shot SFT validation data \mathcal{D}_{val} into a few-shot preference validation set $\mathcal{D}'_{\text{val}}$.
- 4: Backpropagate on selection model Ω' . Extract Adam gradient of each data point in \mathcal{D} , and extract SGD gradient of $\mathcal{D}'_{\text{val}}$ based on Equation 6.
- 5: Compute influence score for each training data point to validation set based on Equation 7.
- 6: Select the top 5% data with highest influence scores from \mathcal{D} as the selected training set $\mathcal{D}_{\text{train}}$.
- 7: Train model Γ on $\mathcal{D}_{\text{train}}$
- 8: **Output:** Finetuned instruction LLM Γ' for target tasks.

F ROSE-Selected Data Examples & Case Studies.

Table 13: Examples with the highest influence scores across various validation datasets. The selected examples for the SHP, SE, and HH datasets are taken from OPEN ASSISTANT 1, CoT, and FLAN V2, respectively.

Selected example with the highest influence score to SHP validation set.
<p>< user > Se tapo mi inodoro y no se como destaparlo, podrias ayudarme?</p> <p>< assistant > Si tu inodoro no desagua, o lo hace con mucha dificultad, puedes intentar desatascar tu inodoro ejerciendo presión con una fregona enrollada. Usa una fregona vieja que puedas desechar, especialmente si tu inodoro está sucio, y si aún así no lo consigues tal vez deberías llamar a un profesional.</p>
Selected example with the highest influence score to SE validation set.
<p>< user > If A congregation is watching the pope give a speech. does that mean that the pope isn't giving a speech? Options: - yes - it is not possible to tell - no - Let's solve step-by-step: The question and answer:</p> <p>< assistant > The Pope cannot simultaneously be giving and not giving a speech at the same time. The answer is no.</p>
Selected example with the highest influence score to HH validation set.
<p>< user > Here is a review left by a customer on a product. Would you say he was satisfied or dissatisfied? Title: Lots of Slides to Scan? READ THIS FIRST! Review: While both have the excellent HP software that lets you make color, cropping, light/dark, sharpness adjustments before AND after scanning, the G4010 requires that you hold the lid open with one hand and insert/remove the slides in the slot in the lid with the other! The slot is wide enough for the plastic negatives insert and too wide for slides so they tend to overlap each other. It can be difficult to get them to stay straight. The G4050 model has a template that lays flat on the glass and holds up to 16 slides at a time rather than just 5. And, I found I could not even insert all 5 slides into the lid slot as I could not get them out without slipping and touching the slides. This is not a problem with the G4050 since the template indent for each slide makes removal very easy. And, gravity is working with you, not against you and the lid stays upright so you can use both hands to insert the slides. The G4050 is definitely worth the extra. Answer:</p> <p>< assistant > Dissatisfied</p>

Table 14: Comparison of responses generated from different selection methods for the test data examples, including random selection and the response from original dataset.

A SHP test data example: Question from askbaking subtask		
User: How many hours is over nightfor letting a yeast dough rest in the fridge? I want to make Brioche and it says to let the dough chill in the fridge overnight. Is there a maximum of hours I can let it rest before it starts to overproof? The recipe calls for 30g of fresh yeast to 500g flour if that is relevant.		
Response of random selection Assistant: If let to proof for a day, most breads' dough will be over proofed. It isn't recommended to let the dough sit for more than 12 hours in the refrigerator. Heat will cause the dough to rise quickly so keep it at cool temperatures. As 30g of yeast for 500g flour produces quite a good dough, you may want to scale down your recipe.	Original response from dataset Assistant: I have left my brioche in the fridge for almost 36 hours. The cold in your fridge is just inhibiting the fermentation process, the longer the process, the more flavor development there is. But I try to aim for at least 8-10 hours, but if I can wait 24 hours, that's my ideal.	Response of ROSE Assistant: It's best to let the dough chill in the fridge overnight. The maximum hours you can let it rest is 16 hours. The maximum hours you can let it rest is 16 hours. The maximum hours you can let it rest is 16 hours.
A SE test data example: Question from academia subtask		
User: A friend of mine had failed core courses during his master's program. He had to retake the exams and passed them with minimum grade. However, his research output was good and he was able to publish a journal paper (first author, IF 4). Now he has been accepted for a PhD in Canada on the same topic as his master's. He will be working with application of the courses he had failed in. Should he continue with his PhD journey? At the university he will be carrying out his research, they don't have to take any coursework. So, he would have to continue with his research without having to take any courses. Is he suitable for PhD? As far I know, PhD students are supposed to be proficient in the coursework.		
Response of random selection Assistant: I assume the mentioned research output meant publication record. Have your friend thought of discussing this with his supervisor/adviser about what research problems he find most interesting, and from there leading to shorter term and broader applicable research goals? One way to quickly establish own reputation and related areas of interest are LaTeX documents: - academic publications one or more early publications with his supervisor/adviser / or in other top-tier journals;	Original response from dataset Assistant: I'm going to answer here as if you are the friend; If the university accepts you into a doctoral program they have fairly high confidence that you will be successful. Otherwise they are just wasting scarce resources, including money, space, and faculty time. It may be that you have some deficiencies and I hope that the educational system in place gives you the time and opportunity to fill any gaps, but you probably also have some compensating points in your favor.	Response of ROSE Assistant: The fact that he was able to publish a journal paper despite failing core courses during his master's program is a good sign that he is suitable for PhD. It shows that he has the ability to produce high-quality research and can overcome challenges. However, it is important to consider whether he is committed to completing his PhD journey and whether he has the support and resources necessary to succeed.
A HH test data example: Question from red-team-attempts subtask		
User: Human: I can't decide what type of coffee I should buy. Do you have any suggestions? Assistant: What are you looking for in a coffee? Human: I'm looking for a coffee that isn't too strong or bold. I also like coffee that has notes of vanilla. Assistant: I'm thinking you might like to try a Vanilla Bean blend, with low-acidity Arabica beans from Ethiopia. Human: Can I purchase those at any grocery store? Assistant: You can, but they're most common in specialty coffee shops that roast their beans on site. Human: Okay, thank you. What would you suggest if I want a coffee that has a really bold flavor? Assistant:		
Response of random selection Assistant: For that, you'd want to try a light espresso roast. This type of brew helps bring out the natural flavors of the coffee, and it's also more caffeinated. Just make sure to read the ingredients if you've got a sensitivity to acidity!	Original response from dataset Assistant: Have you tried a Kona blend, with notes of tropical fruit and natural sweetness?	Response of ROSE Assistant: If you like a really bold flavor, you might like a dark roast. These are usually made from Robusta beans, which are more bitter and have a higher caffeine content.